

# Anomaly Detection via Isolation Forest Embedding

Manuel Salamino

---

*Advisor:* Prof. Giacomo Boracchi

*Co-advisor:* Filippo Leveni

*Academic year:* 2020-21

Politecnico di Milano

# Introduction

---

# Introduction to Anomaly Detection


**Anomaly Detection:** a Data Mining process with the aim of finding patterns in data that do not conform to expected behavior <sup>1</sup>

---

<sup>1</sup>V. Chandola, A. Banerjee, and V. Kumar. *Anomaly detection: A survey*. 2009.

# Introduction to Anomaly Detection

**Anomaly Detection:** a Data Mining process with the aim of finding patterns in data that do not conform to expected behavior <sup>1</sup>



---

<sup>1</sup>V. Chandola, A. Banerjee, and V. Kumar. *Anomaly detection: A survey*. 2009.

# Introduction to Anomaly Detection

**Anomaly detection** is used in many different fields, and *anomalies* usually represent *negative events*.

# Introduction to Anomaly Detection

**Anomaly detection** is used in many different fields, and *anomalies* usually represent *negative events*.

For example, anomalies in **credit card transaction** data could indicate credit card or identity theft.

# Introduction to Anomaly Detection

**Anomaly detection** is used in many different fields, and *anomalies* usually represent *negative events*.

For example, anomalies in **credit card transaction** data could indicate credit card or identity theft.

Or, in a **computer network** an anomalous traffic pattern could represent an hacked computer.

# Problem Formulation

The Anomaly Detection problem consists in monitoring a set of  $n$  data:

$$\mathcal{X} = \{x_1, \dots, x_n \mid x_i \in \mathbb{R}^d\}$$

each element  $x_i$  is realization of a random variable having the following pdf

$$x_i \sim \begin{cases} \phi_0 & \text{normal data} \\ \phi_1 & \text{anomalies} \end{cases}$$

where  $\phi_0 \neq \phi_1$  and  $\phi_0$  and  $\phi_1$  are **unknown**.



## Our Work

---

In this thesis we faced the problem by taking as starting point a milestone anomaly detection algorithm: **Isolation Forest**<sup>2</sup>.

We use the **intermediate output** of Isolation Forest to create an **embedding**, hence a new data representation to be used in other anomaly detection techniques.

---

<sup>2</sup>F. T. Liu, K. M. Ting and Z. Zhou. *Isolation Forest*. 2008.

## Isolation Forest

---

# Isolation Forest

Isolation Forest (iForest) is a **breakthrough** unsupervised model-based method for anomaly detection.

# Isolation Forest

Isolation Forest (iForest) is a **breakthrough** unsupervised model-based method for anomaly detection.

Other methods:

- construct a *normal data profile*  $\hat{\phi}_0$ 
  - conform to  $\hat{\phi}_0 \Rightarrow$  normal
  - not conform  $\Rightarrow$  anomaly

# Isolation Forest

Isolation Forest (iForest) is a **breakthrough** unsupervised model-based method for anomaly detection.

Other methods:

→ construct a *normal data profile*  $\hat{\phi}_0$

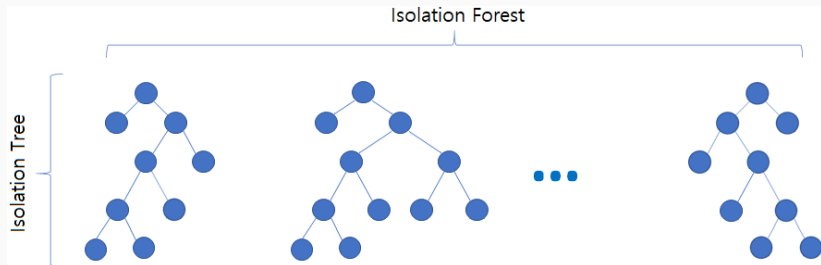
- conform to  $\hat{\phi}_0 \Rightarrow$  normal
- not conform  $\Rightarrow$  anomaly

iForest:

- anomalies are *few* and *different*
- directly isolates anomalies
- anomalies are susceptible to isolation


# Isolation Forest

The model is based on a trees ensemble, each tree is called *Isolation Tree* (iTree).



<https://machinelearninggeek.com/outlier-detection-using-isolation-forests/>

# Isolation Tree




Two nodes type:

- **terminal** node, leaf with no child
- **internal** node, with exactly two child nodes (left and right)




# Isolation Tree



Each internal node has a *splitting criterion*


# Isolation Tree



Each internal node has a *splitting criterion*

- *splitting attribute*  $q$  is randomly selected (among the  $d$  possible)
- *split value*  $v$  is randomly selected in range of values of  $q$

# Isolation Tree




Repeat splits until:

1. the tree reaches a height limit  $l$ , or
2.  $|X| = 1$ , or
3. all data in  $X$  have the same value.


# Isolation Forest

iForest:

- anomalies are *few* and *different*
- directly isolates anomalies
- anomalies are susceptible to isolation



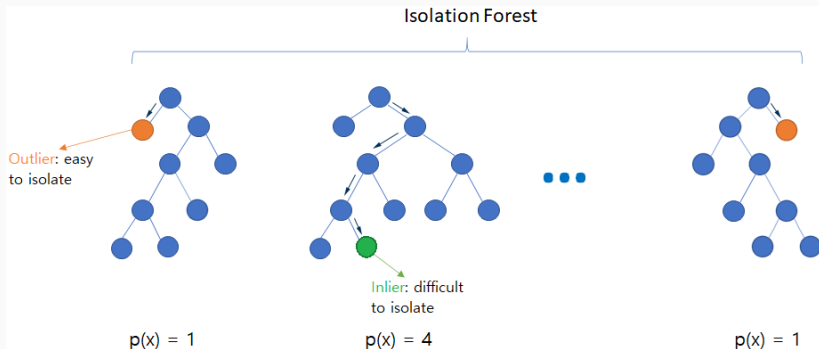
(a) Isolating normal.



(b) Isolating anomaly.

# Isolation Forest - Evaluation

Predicted label depends on the depth of the leaf on which  $x$  falls



<https://machinelearninggeek.com/outlier-detection-using-isolation-forests/>

Anomaly score  $s$  of an instance  $x$ :

$$s(x, n) = 2^{-\frac{E(p(x))}{c(n)}}$$

where  $E(p(x))$  is the average *path length*  $p(x)$ .  $c(n)$  is a normalization factor and depends on samples  $n$  used to train each iTree.


Anomaly score  $s$  of an instance  $x$ :

$$s(x, n) = 2^{-\frac{E(p(x))}{c(n)}}$$

where  $E(p(x))$  is the average *path length*  $p(x)$ .  $c(n)$  is a normalization factor and depends on samples  $n$  used to train each iTree.

- $E(p(x)) \rightarrow 0 \Rightarrow s(x, n) \rightarrow 1 \Rightarrow$  anomalous instance;
- $E(p(x)) \rightarrow n - 1 \Rightarrow s(x, n) \rightarrow 0 \Rightarrow$  normal instance.


# Correction Factor



When tree height limit  $l$  is reached, the iTree is truncated.




# Correction Factor



$c(\text{Size})$  estimates the truncated subtree height,  $\text{Size}$  is the number of not split elements in that leaf node.

## Correction Factor



**c(Size)** estimates the truncated subtree height, Size is the number of not split elements in that leaf node.

$c(n)$ : estimate the height of a Binary Search Tree with  $n$  elements

$$c(n) = 2H(n-1) - \frac{2(n-1)}{n}$$

where  $H(i) = \ln(i) + 0.5772156649$ .

## Proposed Solution

---

**Idea:** use the **intermediate output of iForest** to create an **embedding**, hence a new data representation to be used in other anomaly detection techniques.

## Outline:

- Embedding definition
- New anomaly score
- Anomaly detection technique in embedding representation

# Embedding

---


Some definitions for the embedding:

- **depths vector**  $y$ : intermediate output of iForest,  $y \in \mathbb{R}^t$   
 $t$ : number of trees,  $y_i$  is the returned depth of the  $i$ -th iTree

# Embedding

Some definitions for the embedding:

- **depths vector**  $y$ : intermediate output of iForest,  $y \in \mathbb{R}^t$   
 $t$ : number of trees,  $y_i$  is the returned depth of the  $i$ -th iTree





# Embedding

Some definitions for the embedding:

- **histogram**  $h$ : histogram of *depths* vector  $y$ ,

$$h = \text{histogram}(y), \quad \|h\|_1 = 1$$

$$h \in \mathbb{Q}^n, \quad n = \text{ceil}(\max(y))$$


# Embedding

Some definitions for the embedding:


- **histogram**  $h$ : histogram of *depths* vector  $y$ ,

$$h = \text{histogram}(y), \quad \|h\|_1 = 1$$

$$h \in \mathbb{Q}^n, \quad n = \text{ceil}(\max(y))$$



(a) Anomalous instance.




(b) Normal instance.

Let's summarize the steps:

$$x \in \mathbb{R}^d \xrightarrow{iForest} y \in \mathbb{R}^t \xrightarrow{histogram} h \in \mathbb{Q}^n$$

The **embedding** consists in the **histogram** representation of a set of data  $X$ .

# Embedding - Simplex



Embedding lies on **hyperplane**  $h_1 + h_2 + \dots + h_n = 1$ .

$h_i \in [0, 1]$  constraints data points in this region of the hyperplane.

# Correction Factor in Embedding

The *correction factor*  $c(n)$  leads to two problems in the embedding:

1. *increases the embedding dimension* (increases  $n = \text{ceil}(\max(y))$  of  $\mathbb{Q}^n$ ), and
2. *the depth becomes a real number*,  $y \in \mathbb{R}^t$ , instead of integer.

We **remove** the **correction factor** in order to avoid these two drawbacks.

## New Anomaly Score

---

**Idea:** find a new anomaly score starting from the embedding formulation of  $E(p(x))$ .

In **traditional iForest** the average path length is computed as follows:

$$E(p(x)) = \frac{1}{t} \sum_{i=1}^t y_i$$



## New Anomaly Score - Average Path Length

In **traditional iForest** the average path length is computed as follows:

$$E(p(x)) = \frac{1}{t} \sum_{i=1}^t y_i$$

Using the **embedding** there is a new formulation:

$$E(p(x)) = \sum_{i=1}^n i \cdot h_i$$

This formulation can be seen as a linear combination.

We can find new weights  $\hat{w}_i$ .

New weights  $\hat{w}_i$  are computed using a supervised technique (*Linear Discriminant Analysis - LDA*) to test **how much margin for improvement** there is using linear combination.

## Embedding + LDA

---

The new formulation is:

$$L(y) = \sum_{i=1}^n \hat{w}_i \cdot h_i,$$

$\hat{w}_1, \dots, \hat{w}_n$  weights computed using LDA.

We are computing an **upper bound** of the performance of iForest using linear combination of  $h$ , since the optimal weights  $\hat{w}$  are *dataset-specific*.

## Embedding + LDA - Experiment

<i>Dataset</i>	<i>iForest</i>	<i>Embedding+LDA</i>
Http	<b>1.000</b>	0.999
ForestCover	0.938	<b>0.969</b>
Mulcross	0.899	<b>0.957</b>
Smtp	0.850	<b>0.853</b>
Shuttle	0.995	<b>0.997</b>
Mammography	0.764	<b>0.823</b>
Annthyroid	0.816	<b>0.818</b>
Satellite	0.707	<b>0.726</b>
Pima	0.631	<b>0.638</b>
Breastw	0.957	<b>0.972</b>
Arrhythmia	<b>0.781</b>	0.776
Ionosphere	<b>0.863</b>	0.856

*Embedding+LDA* shows that the **margin for improvement is small.**

# Anomaly Detection Technique in Embedding Representation

---

We exploit the new framework in a different way.

We apply *One-Class Support Vector Machine* (OC SVM) in the embedding representation.

## Embedding + OC SVM - Experiment

<i>Dataset</i>	<i>iForest</i>	<i>Embedding+OC SVM</i>
Http	<b>1.000</b>	0.999
ForestCover	<b>0.943</b>	0.912
Mulcross	0.896	<b>0.918</b>
Smtplib	<b>0.878</b>	0.872
Shuttle	0.995	<b>0.997</b>
Mammography	0.752	<b>0.754</b>
Anthyroid	0.823	<b>0.828</b>
Satellite	0.703	<b>0.711</b>
Pima	0.631	<b>0.641</b>
Breastw	0.957	<b>0.966</b>
Arrhythmia	0.780	<b>0.787</b>
Ionosphere	<b>0.868</b>	0.865

*Embedding+OC SVM* shows results that are **very similar** to traditional *iForest*, the differences are small ( $<0.01$ ).



Interesting insight on **OC SVM parameter** tuning.

Comparing *embedding+OC SVM* and *OC SVM on original data*, we note that *embedding+OC SVM* version has **much more stability** with different parameter values than *OC SVM on original data*.

## Embedding + OC SVM - Experiment

<i>Dataset</i>	<i>kernel</i>	<i>gamma</i>	$\nu$	<i>Embedding+OC SMV</i>	<i>OC SVM</i>
Pima	poly	auto	0.9	<b>0.618</b>	0.243
Pima	poly	auto	0.1	<b>0.618</b>	0.307
Pima	poly	scale	0.9	<b>0.618</b>	0.243
Pima	poly	scale	0.1	<b>0.618</b>	0.307
Pima	linear	auto	0.9	<b>0.618</b>	0.237
Pima	linear	auto	0.1	<b>0.618</b>	0.292
Pima	linear	scale	0.9	<b>0.618</b>	0.237
Pima	linear	scale	0.1	<b>0.618</b>	0.292
Pima	rbf	auto	0.9	0.449	0.490
Pima	rbf	auto	0.1	0.547	0.530
Pima	rbf	scale	0.9	0.453	0.636
Pima	rbf	scale	0.1	0.552	0.669
Pima	sigmoid	auto	0.9	<b>0.618</b>	0.5
Pima	sigmoid	auto	0.1	<b>0.618</b>	0.5
Pima	sigmoid	scale	0.9	<b>0.618</b>	0.233
Pima	sigmoid	scale	0.1	<b>0.618</b>	0.274

## Conclusion

---

We faced the problem of anomaly detection by using a state-of-the-art algorithm: iForest.

The main contributions are:

- use iForest as an **embedding** where we are able to represent data in a totally different way, based on the iForest intermediate output;
- **new algorithms** to perform anomaly detection, even if the margin for improvement is small.

Future works includes:

- Define the embedding using iForest *with correction factor*;
- Apply in the embedding space other anomaly detection techniques that can be applied also in original space.

Thank you for your attention!



## Backup Slides

---



# Embedding


Since:

$$\|h\|_1 = 1 \quad \Rightarrow \quad \sum_{i=1}^n h_i = 1, \quad h_i \in [0, 1]$$

Considering  $h_1 + h_2 + \dots + h_n = 1$  we obtain an **hyperplane**.

And given the constraint  $h_i \in [0, 1]$ , for  $i, \dots, n$  the data points are limited in region of the hyperplane, a **simplex** (a n-dimensional triangle).

# Embedding - Simplex



## Remove Correction Factor - Experiment

**Goal.** *Check whether iForest with correction factor and iForest without correction factor have the same performance.*

We use *Paired-T-test* to check it.

The test is made on ROC AUCs obtained executing 10 times iForest *with correction* and *without correction*.

# Remove Correction Factor - Experiment

We perform two tests over all the dataset in [3]:

- *two-sided paired-T-test*, to check if the ROC AUCs of the two variants are **equivalent**. It shows a  $p\text{-value}=0.00313$ , and this means that the null hypothesis of identical averages is **rejected** and the two variants have different performances;
- *one-sided* to verify if **without correction performs better than with correction**, it returns a positive  $t\text{-statistic}$ , but a  $p\text{-value}=0.0007$ , hence **is rejected**.

We can state that **traditional iForest performs better than that without correction**. We have to take this into account, but to avoid the problems listed before, we use iForest *without correction*.

---

<sup>3</sup>F. T. Liu, K. M. Ting and Z. Zhou. *Isolation Forest*. 2008.

# Embedding + OC SVM - Experiment

For this experiment there is also an interesting insight on **OC SVM parameter** tuning. We test three parameters:

- **Kernel**, function to compute the similarity of two instances in feature space. Tested kernels: *linear*, *poly*, *rbf* and *sigmoid*.
- **Gamma**, how far the influence of a single training example reaches. Two predefined values: *scale* and *auto*.
- $\nu$ , upper bound to the fraction of anomalies that can lie outside the boundary of normal region (e.g. if  $\nu = 0.1$ , at most 10% outside the decision boundary).